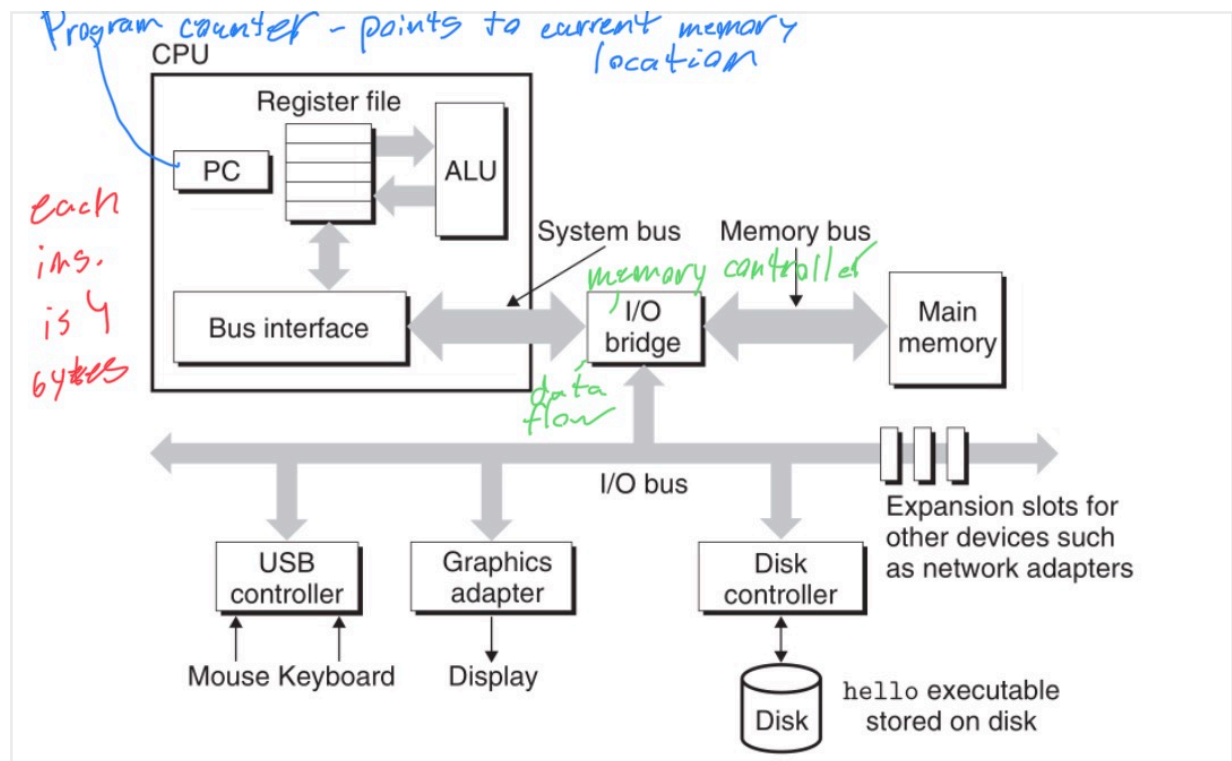


Slides 1



MIPS Registers

Register Number	Conventional Name	Usage
\$0	\$zero	Hard-wired to 0
\$1	\$at	Reserved for pseudo-instructions
\$2 - \$3	\$v0, \$v1	Return values from functions
\$4 - \$7	\$a0 - \$a3	Arguments to functions - not preserved by subprograms
\$8 - \$15	\$t0 - \$t7	Temporary data, not preserved by subprograms
\$16 - \$23	\$s0 - \$s7	Saved registers, preserved by subprograms
\$24 - \$25	\$t8 - \$t9	More temporary registers, not preserved by subprograms
\$26 - \$27	\$k0 - \$k1	Reserved for kernel. Do not use.
\$28	\$gp	Global Area Pointer (base of global data segment)
\$29	\$sp	Stack Pointer
\$30	\$fp	Frame Pointer
\$31	\$ra	Return Address
\$f0 - \$f3	-	Floating point return values
\$f4 - \$f10	-	Temporary registers, not preserved by subprograms
\$f12 - \$f14	-	First two arguments to subprograms, not preserved by subprograms
\$f16 - \$f18	-	More temporary registers, not preserved by subprograms
\$f20 - \$f31	-	Saved registers, preserved by subprograms

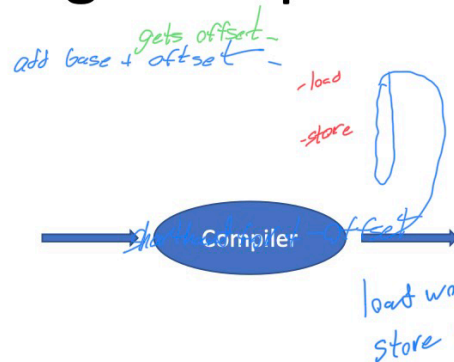
High-level Language

- Implementing a swap function in C

```

swap(int v[], int k) {
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}

```



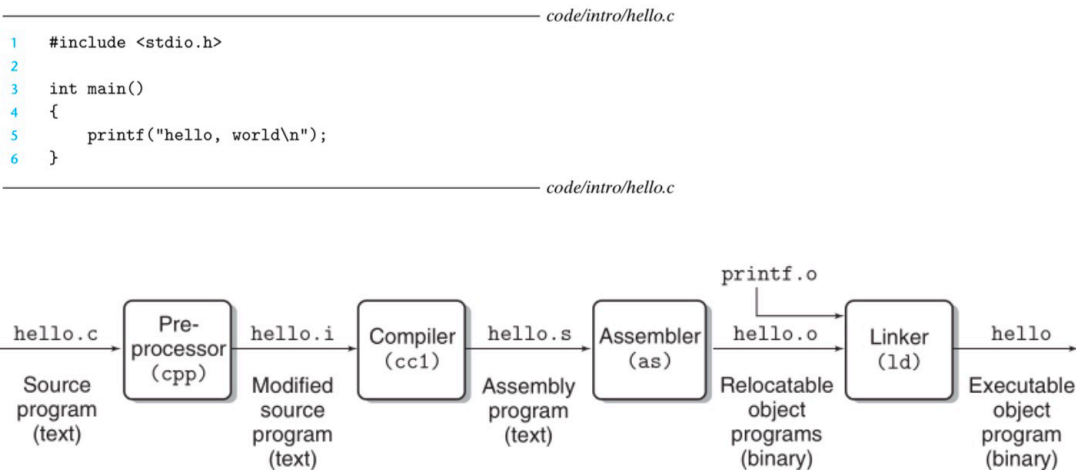
swap:

```

multi $2, $5, 4
add   $2, $4, $2
lw    $15, 0($2)
lw    $16, 4($2)
sw    $16, 0($2)
sw    $15, 4($2)
jr    $31

```

Executable Object Program from a Source Program



Microarchitecture: implementation to support ISA

ISA: the interface.

- **Word**
 - **Default** data size for computation
 - 8-bit, 16-bit, 32-bit, and 64-bit processor
- **Address**
 - Points to a location in memory in byte (byte addressable)
 - 32-bit address, 2^{32} bytes = 4GB
 - 28-bit address required if you have 256MB memory
- **Cache**
 - Faster but smaller -> expensive

Evaluating Performance

- What does “fast” mean?
 - Speed VS Throughput (Bandwidth)
 - CPI (cycles per instruction) as speed
 - IPC (Instructions per cycles) as throughput

Airplane	Passenger capacity	Cruising range (miles)	Cruising speed (m.p.h.)	Passenger throughput (passengers x m.p.h.)
Boeing 777	375	4630	610	228,750
Boeing 747	470	4150	610	286,700
BAC/Sud Concorde	132	4000	1350	178,200
Douglas DC-8-50	146	8720	544	79,424

Improve throughput with: pipelining, hyper threading, clock speed

Evaluating Performance

- Instructions per second
 - E.g., 3GHz with Intel i7 (Quad Cores) and IPC = 1
 - 3 billion cycles per second
 - 4 issues-superscalars
 - Hyper-Threading (fetching two programs every cycle)
 - $3 \times 10^9 \times 1 \times 4 \times 4 \times 2$ / second
 - = 96 billion instructions / second at the maximum speed

Different ISA's effect actual perf.



- Execution time per program
 - $NI * CPI * \text{clock_cycle_time}$
 - Where, NI: number of Instructions – Small is better
 - CPI (clock cycle per instruction) – Small is better
 - Clock cycle time – Small is better (Higher clock is better)

Memory Performance

System Event	Actual Latency	Scaled Latency
One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	~100 ns	4 min
Intel® Optane™ DC persistent memory access	~350 ns	15 min
Intel® Optane™ DC SSD I/O	<10 μ s	7 hrs
NVMe SSD I/O	~25 μ s	17 hrs
SSD I/O	50–150 μ s	1.5–4 days
Rotational disk I/O	1–10 ms	1–9 months
Internet call: San Francisco to New York City	65 ms ^[3]	5 years
Internet call: San Francisco to Hong Kong	141 ms ^[3]	11 years

4 ghz dual core
3-way scalar, hyper threaded

